Table 1: Clone-generation results on the industrial measurement and control corpus. Twelve programs × three seeds = 36 runs per configuration. *Equiv. conf.* is reported over categories C and D only (54 runs), the subjects that ship golden vectors and invariants; categories A and B have no equivalence probes and are excluded rather than scored as passes.

Configuration	Gen. success	Equiv. conf.	Divergence	Repair iters	Cost (USD)
Single-shot LLM (B1)	28/36	15/18	0.250	0 (0 re-entry)	$0.79
Proteus w/o verification (B2)	32/36	16/18	0.135	0 (0 re-entry)	$4.74
Proteus (full)	26/36	13/18	0.175	6 (6 re-entry)	$6.12

Table 2: Blind LLM review of all 108 generated clones. An independent model (`claude-opus-5`, distinct from the `claude-sonnet-5` used for generation) was shown each original and its clone with all system-identifying markers removed, and asked for concrete behavioural defects. NO-OP denotes a clone byte-identical to its original: trivially faithful because nothing was attempted, and therefore excluded from the clean rate rather than counted as a success.

Configuration	BROKEN	MAJOR	MINOR	CLEAN	NO-OP	Clean rate
Single-shot LLM (B1)	4	21	8	3	0	3/36
Proteus w/o verification (B2)	2	8	1	14	11	14/25
Proteus (full)	6	10	4	13	3	13/33

Table 3: Precision of each configuration's own success signal. Of the clones a configuration certified (compiles, passes the full suite, and genuinely substituted the dependency), the fraction an independent blind reviewer found to carry BROKEN or MAJOR behavioural defects. Accepting more clones and accepting sound clones are different capabilities, and they trade off against one another.

Configuration	Clones accepted	With serious defects	Defect rate
Single-shot LLM (B1)	21/36	17	**81%**
Proteus w/o verification (B2)	14/36	6	**43%**
Proteus (full)	11/36	3	**27%**